


MAE Learn from what is missing

1 Mask the image



Original image

16 patches

75% at random



4 visible · 12 hidden

One synthetic image; no model run.
Projection layers and CLS token omitted.

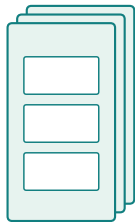
2 Encode

Pack visible patches



4 tokens + pos.

Encoder



visible only

3 Restore and predict

Shared mask token



× 12 copies



Unshuffle, then add decoder positions.

Original positions



4 latents + 12 masks
Original targets

Decoder



lightweight

Pixel predictions

16 symbolic slots



12 masked slots



MSE

masked only

