


MAE: visible-only encoding during pretraining

The decoder reconstructs masked patches; recognition reuses the encoder.



Dashed: loss supervision. Masked-patch loss only. Patch projection and positional embeddings omitted.