


MAE: encode fewer patches, retain recognition quality

ViT-L/16 · ImageNet-1K · 800 epochs · 128 TPU-v3 cores

a Learn from visible patches

Illustrative image patch grid



↓ Keep 25%
for the encoder



**Reconstruct what
is missing**

Mask tokens enter the decoder later.

b Quality and training time

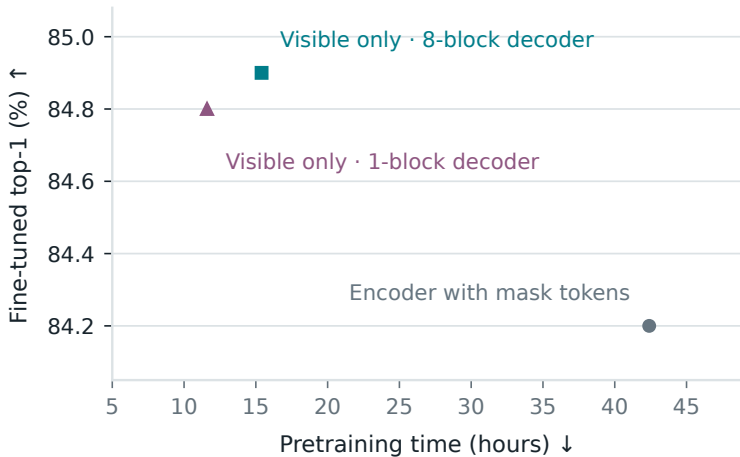


Table 2 ablation: 63.7% less time and +0.7 percentage points for the 8-block variant.