



MAE Learn from visible content; reconstruct what is missing

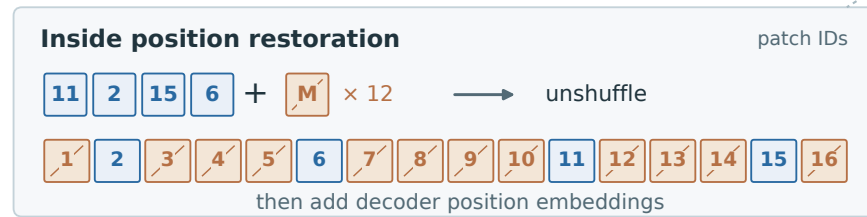
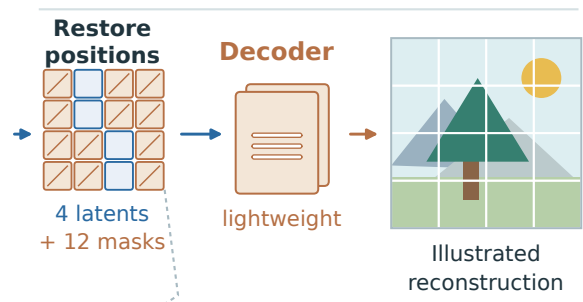
a Randomly keep 25%



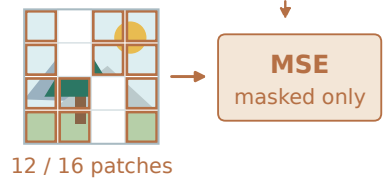
b Encode



c Reconstruct



Original targets



Scene and reconstruction are illustrative; 16-patch example; no model run.
Pretraining only. CLS token, projection layers and internal attention details omitted.

MAE v2 | 7-inch-wide draft | illustrative vector scene