

Reasoning Gym

Generate a problem; verify the answer

Procedural benchmark setup · one selected leg-counting example

a Configure

Task generator

leg_counting

Seed 42

Size 10

Params defaults

 10 entries

b Generate one entry

TEXT QUESTION

How many legs in total?

 1 sea slug

 1 deer

Only the question goes to the model

HELD BY EVALUATOR

Reference answer: 4

Metadata: animal counts

c Answer and verify

 **Reasoning model**

Candidate: **4**

Task verifier

score_answer

1.0
score

Schematic: drawings explain the text prompt; candidate 4 is illustrative (documented score: 1.0).

Other tasks use their own verifiers and may accept multiple valid solutions.