

SEAL Train the editor through the results of adaptation

(a) Training: candidate edits are judged after updating the model

1 Generate

CONTEXT C
Passage or examples

θ_t
LM
Current policy

Next policy θ_{t+1}

EDIT 1

EDIT 2

EDIT 3

2 Adapt each copy

All copies start from the same θ_t

Apply edit
+ SFT

$\theta'_{t,1}$

LM

Apply edit
+ SFT

$\theta'_{t,2}$

LM

Apply edit
+ SFT

$\theta'_{t,3}$

LM

Inner loop: gradient-based adaptation

3 Test & reward

Held-out task τ

Task score



Task score



Task score



Rewards shown schematically

4 Reinforce useful edits

SFT on selected context-edit pairs

RL RECORD

C, edit_i, r_i

(b) What is a self-edit? Two domain-specific representations


KNOWLEDGE → TRAINING TEXT

Passage: "Planet Lyra orbits two suns."

Edit: "Lyra is a planet. It has two suns."

FEW-SHOT TASK → TRAINING CONFIG

Augmentation



Controls training:

augment: rotations, flips

choose: learning rate, epochs

Illustrative passage and grids; schematic rewards. The source paper defines domain-specific selection.