


Transformer: attention connects source and target sequences

Original post-norm blocks: $\text{LayerNorm}(x + \text{Sublayer}(x))$. Residual wiring omitted in this overview.



Schematic. Embeddings and prediction head sit outside repeated blocks. Token feedback omitted.